

Data Exploration and Preprocessing:- Analyze the dataset for trends, missing values, and outliers. Perform data cleaning, feature engineering, and transformations.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split

# Load Dataset
df = pd.read_csv(r"C:\Users\Pravallika\Downloads\Airbnbdata.csv")

# Quick Overview
print("Dataset Shape:", df.shape)
print("\nDataset Info:")
print(df.info())
print("\nMissing Values:\n", df.isnull().sum())
```

Dataset Shape: (74111, 29)

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74111 entries, 0 to 74110
Data columns (total 29 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               74111 non-null   int64  
 1   log_price        74111 non-null   float64 
 2   property_type    74111 non-null   object  
 3   room_type        74111 non-null   object  
 4   amenities        74111 non-null   object  
 5   accommodates     74111 non-null   int64  
 6   bathrooms         73911 non-null   float64 
 7   bed_type          74111 non-null   object  
 8   cancellation_policy 74111 non-null   object  
 9   cleaning_fee     74111 non-null   bool    
 10  city              74111 non-null   object  
 11  description       74111 non-null   object  
 12  first_review      58247 non-null   object  
 13  host_has_profile_pic 73923 non-null   object  
 14  host_identity_verified 73923 non-null   object  
 15  host_response_rate 55812 non-null   object  
 16  host_since        73923 non-null   object  
 17  instant_bookable  74111 non-null   object  
 18  last_review        58284 non-null   object  
 19  latitude           74111 non-null   float64 
 20  longitude          74111 non-null   float64 
 21  name               74111 non-null   object  
 22  neighbourhood      67239 non-null   object
```

```
23 number_of_reviews      74111 non-null  int64
24 review_scores_rating   57389 non-null  float64
25 thumbnail_url          65895 non-null  object
26 zipcode                 73143 non-null  object
27 bedrooms                74020 non-null  float64
28 beds                    73980 non-null  float64
dtypes: bool(1), float64(7), int64(3), object(18)
memory usage: 15.9+ MB
None
```

Missing Values:

id	0
log_price	0
property_type	0
room_type	0
amenities	0
accommodates	0
bathrooms	200
bed_type	0
cancellation_policy	0
cleaning_fee	0
city	0
description	0
first_review	15864
host_has_profile_pic	188
host_identity_verified	188
host_response_rate	18299
host_since	188
instant_bookable	0
last_review	15827
latitude	0
longitude	0
name	0
neighbourhood	6872
number_of_reviews	0
review_scores_rating	16722
thumbnail_url	8216
zipcode	968
bedrooms	91
beds	131

dtype: int64

```
# Analyze Trends
# Price Distribution
```

```
plt.figure(figsize=(8,5))
sns.histplot(df['log_price'], bins=50, kde=True, color='skyblue')
plt.title("Distribution of Airbnb Prices")
plt.xlabel("log_Price")
plt.ylabel("Count")
```

```

plt.show()

# Average Price by Room Type
if 'room_type' in df.columns:
    plt.figure(figsize=(7,5))
    sns.barplot(x='room_type', y='log_price', data=df,
estimator='mean', errorbar=None, palette='viridis')
    plt.title("Average Price by Room Type")
    plt.show()


# Correlation Heatmap
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm',
fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()

# Check for Outliers

plt.figure(figsize=(8,5))
sns.boxplot(x=df['log_price'], color='salmon')
plt.title("Price Outliers")
plt.show()


Q1 = df['log_price'].quantile(0.25)
Q3 = df['log_price'].quantile(0.75)
IQR = Q3 - Q1
upper = Q3 + 1.5 * IQR
lower = Q1 - 1.5 * IQR
print(f"Outlier Threshold: {lower:.2f} to {upper:.2f}")


```




```
C:\Users\Pravallika\AppData\Local\Temp\  
ipykernel_23216\1686290662.py:14: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be  
  removed in v0.14.0. Assign the `x` variable to `hue` and set  
  `legend=False` for the same effect.  
  
  sns.barplot(x='room_type', y='log_price', data=df, estimator='mean',  
  errorbar=None, palette='viridis')
```

Average Price by Room Type





Price Outliers



```
Outlier Threshold: 2.96 to 6.57
```

```
# Handle Missing Values
df['review_scores_rating'] = df['review_scores_rating'].fillna(0)
df = df.dropna(subset=['log_price'])

# Remove Outliers
df = df[(df['log_price'] >= lower) & (df['log_price'] <= upper)]
print("Data shape after removing outliers:", df.shape)

Data shape after removing outliers: (72579, 29)

# Feature Engineering

if 'amenities' in df.columns:
    df['num_amenities'] = df['amenities'].apply(lambda x:
len(str(x).split(',')))

# Encode categorical variables

categorical_cols = df.select_dtypes(include='object').columns
le = LabelEncoder()
```

```

for col in categorical_cols:
    df[col] = le.fit_transform(df[col].astype(str))

# Feature Transformation

scaler = StandardScaler()
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
df[numeric_cols] = scaler.fit_transform(df[numeric_cols])

print("Data Preprocessing Complete")

```

□ Data Preprocessing Complete

Model Development:- Build a regression model to predict listing prices.

```

# Split Data into Train, Validation, and Test Sets

X = df.drop('log_price', axis=1)
y = df['log_price']

X_train, X_temp, y_train, y_temp = train_test_split(X, y,
test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp,
test_size=0.5, random_state=42)

print(f"Train: {X_train.shape}, Validation: {X_val.shape}, Test: {X_test.shape}")

# Impute missing values using the mean of the training set'

for col in X_train.columns:
    if X_train[col].isnull().any():
        mean_val = X_train[col].mean()
        X_train[col] = X_train[col].fillna(mean_val)
        X_val[col] = X_val[col].fillna(mean_val)
        X_test[col] = X_test[col].fillna(mean_val)

# Train Multiple Regression Models

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor

lr = LinearRegression()
rf = RandomForestRegressor(random_state=42)
xgb = XGBRegressor(random_state=42)

lr.fit(X_train, y_train)
rf.fit(X_train, y_train)

```

```

xgb.fit(X_train, y_train)

print("Model Training Completed")

Train: (50805, 29), Validation: (10887, 29), Test: (10887, 29)
Model Training Completed

```

Model Evaluation Evaluate the model's performance using appropriate metrics like RMSE, MAE, and R².

```

from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

def evaluate_model(model, X, y):
    preds = model.predict(X)
    rmse = np.sqrt(mean_squared_error(y, preds))
    mae = mean_absolute_error(y, preds)
    r2 = r2_score(y, preds)
    return rmse, mae, r2

models = {'Linear Regression': lr, 'Random Forest': rf, 'XGBoost': xgb}

print("\nModel Evaluation Results:")
for name, model in models.items():
    rmse, mae, r2 = evaluate_model(model, X_test, y_test)
    print(f"\n{name}:")
    print(f"RMSE: {rmse:.3f}")
    print(f"MAE: {mae:.3f}")
    print(f"R2: {r2:.3f}")

# Visual comparison

results = pd.DataFrame({
    'Model': ['Linear Regression', 'Random Forest', 'XGBoost'],
    'RMSE': [evaluate_model(lr, X_test, y_test)[0],
              evaluate_model(rf, X_test, y_test)[0],
              evaluate_model(xgb, X_test, y_test)[0]],
    'R2 Score': [evaluate_model(lr, X_test, y_test)[2],
                   evaluate_model(rf, X_test, y_test)[2],
                   evaluate_model(xgb, X_test, y_test)[2]]
})

plt.figure(figsize=(8,5))
sns.barplot(x='Model', y='R2 Score', data=results, palette='magma')
plt.title("Model Comparison by R2 Score")
plt.show()

```

Model Evaluation Results:

```
Linear Regression:
```

```
RMSE: 0.692
```

```
MAE: 0.537
```

```
R2: 0.524
```

```
Random Forest:
```

```
RMSE: 0.555
```

```
MAE: 0.417
```

```
R2: 0.694
```

```
XGBoost:
```

```
RMSE: 0.542
```

```
MAE: 0.408
```

```
R2: 0.708
```

```
C:\Users\Pravallika\AppData\Local\Temp\  
ipykernel_23216\834269041.py:32: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be  
removed in v0.14.0. Assign the `x` variable to `hue` and set  
'legend=False' for the same effect.
```

```
sns.barplot(x='Model', y='R2 Score', data=results, palette='magma')
```

